Hollow House Institute | Policy Brief

The Slop Problem

Executive Summary — AI-Generated Pollution and the Case for Execution-Time Governance

Amy | Founder, Hollow House Institute May 2026

A condensed briefing for board members, policymakers, and senior leadership.

AI-generated content is flooding every major information channel — research databases, news aggregators, policy repositories, social platforms, and academic preprint servers. The term for this phenomenon is "slop": content generated at machine speed, without meaningful human review, accountability, or provenance. It is not a fringe problem. It is structural. And left unaddressed, it will degrade the information environments that governance, research, and public trust depend on.

Why This Is Happening Now

Three forces have converged to produce this crisis.

First, the cost of generating content has collapsed to near-zero. A single person with API access can produce tens of thousands of articles, abstracts, or reports per day. Volume is no longer a resource constraint — it is the default.

Second, platforms reward volume. The engagement and indexing algorithms that determine what gets seen do not distinguish quality from quantity. Content that is produced fast and published often rises. Content that is carefully verified does not receive a corresponding advantage.

Third, bad actors have learned that AI-normalized content provides cover. When slop is everywhere, targeted disinformation becomes significantly harder to identify. The noise floor rises, and signal becomes harder to locate.

There is a fourth dimension that makes this a compounding problem, not merely a chronic one. Peer-reviewed research published in *Nature* (Shumailov et al., 2024) established that when AI-generated content is scraped back into AI training data, it causes cumulative and irreversible degradation of future model quality — a phenomenon termed "model collapse." The AI systems of tomorrow are being trained on the slop of today. This is not theoretical. It is already occurring.

What Is at Stake

The risks are concrete and immediate across three critical domains.

Research integrity. Fake and AI-fabricated abstracts, citations, and summaries are circulating inside academic and policy research pipelines. Scholars and analysts are citing content that was never independently verified — and in some cases, content that never existed as a primary source at all. The downstream effect is a corruption of the knowledge base that policy decisions rely upon.

Policy quality. Governments and institutions that are making AI-informed decisions are increasingly downstream of AI-generated summaries, not primary sources. The distance between a policy choice and a verifiable fact is growing. That gap is where errors compound and accountability disappears.

Epistemic trust. When the information environment cannot be trusted, people default to instinct, ideology, or whoever shouts loudest. That is not merely a media problem. It is a governance failure with direct institutional consequences — for public health, democratic legitimacy, regulatory effectiveness, and crisis response.

> ***"The governance window is open now. The longer it remains unaddressed, the more entrenched the incentive stack becomes."***

The Missing Layer: Execution-Time Governance

Almost all current AI regulation — including the EU AI Act, executive orders, and national AI strategies — focuses on model training and deployment licensing. These are necessary interventions. But none of them intercept the moment when AI output is actually generated and pushed into the information environment. That moment — execution time — is where slop is produced. And it is currently ungoverned.

HHI's position is that governance must operate at the point of execution, not only at the point of design. The design-phase frame gives regulators visibility into what models are capable of. It does not give them visibility into what those models are actually doing, at scale, in real time.

In practice, execution-time governance means several things: human review checkpoints before AI content is published in high-stakes contexts; machine-readable provenance — a tamper-evident record of how content was created — embedded at generation time rather than appended later; mandatory human attestation for research, news, and policy content, with legal liability attached to false claims; and auditable logs from AI content generation APIs that are accessible to regulators, not only to the companies that operate them.

These mechanisms are not aspirational. HHI operates a working governance stack — including append-only telemetry, cryptographic audit chains, human stop authority, and continuous behavioral drift monitoring — that demonstrates these capabilities are buildable and operable today, with existing technology.

What Governance Needs to Do — Five Priorities

1. **1.**
2. **Establish mandatory AI content provenance standards** at the platform and API level — machine-readable, tamper-evident, and embedded at generation time, not added later as an optional metadata field.
3. **2.**
4. **Require human attestation** for content published in research, news, and policy contexts — and attach meaningful legal liability to false attestation, so that accountability is real rather than procedural.
5. **3.**
6. **Implement execution-time governance frameworks** in high-stakes publishing pipelines, drawing on and extending existing standards including NIST AI RMF 1.0, ISO/IEC 42001, and the human oversight requirements of the EU AI Act (Articles 13 and 14).
7. **4.**
8. **Establish platform accountability** — platforms that algorithmically amplify AI-generated slop should bear partial liability for downstream epistemic harms, on the same logic that product liability applies to manufacturers whose distribution decisions cause harm.
9. **5.**
10. **Fund independent provenance infrastructure** — not owned by AI companies or platforms — and invest in information literacy at every level of public education, so that the human layer of the information ecosystem becomes more resilient, not less.

HHI's Position

Hollow House Institute is a focused, independent research and policy institute based in Arlington, Texas. This paper is part of our ongoing work to name what is happening clearly, and to build frameworks that can be operationalized by institutions that want to act — not just study the problem. We are not neutral observers. We believe the information environment is a public good, and that its integrity is worth defending.

The governance window is open now. The longer it remains unaddressed, the more entrenched the incentive stack becomes — for platforms, for API providers, and for bad actors who depend on the current absence of accountability. We are available to engage directly with policymakers, institutional leaders, standards bodies, and boards who are ready to move from concern to action.

Amy

Founder, Hollow House Institute

Arlington, Texas | May 2026

hollow.house

For the full white paper, bibliography, standards crosswalk, and runtime evidence appendix, see:

"The Slop Problem: AI-Generated Pollution and the Case for Execution-Time Governance"

Hollow House Institute, May 2026 | hollow.house